

Neural Networks

Jordi Villà i Freixa

Universitat de Vic - Universitat Central de Catalunya
Study Abroad

jordi.villa@uvic.cat

course 2023-2024

1 Introduction to NN

2 Bibliography

What to expect?

In this session we will discuss:

- Neural networks
- Gradient Descent
- Backpropagation

Evolution of AI

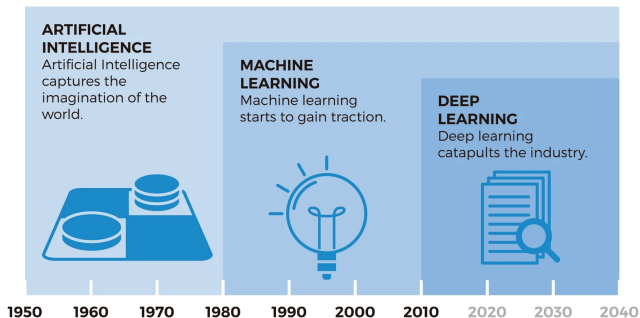


Figure 1: Evolution of AI, ML and DL. Source [BISMART](#).

Supervised ML

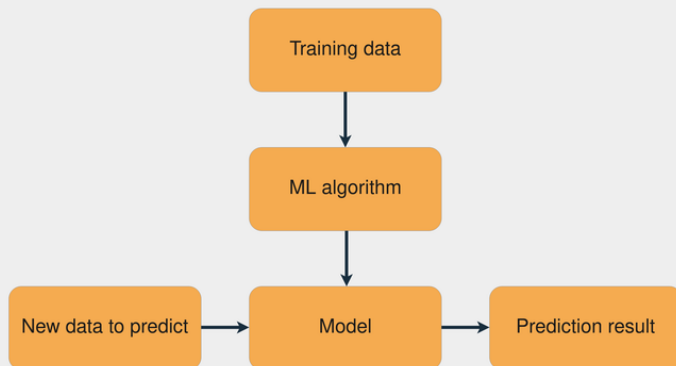


Figure 2: Workflow to train a model using supervised learning.

Supervised ML

- Build models capable of learning from a series of data (x)
- The learning process depends on the task we want to train the model:

Supervised the input data X is accompanied by the values y that we want the model to learn or also called targets (linear regression, decision trees, ...). The objective is that the prediction of the model given some data x is equal to the target.

Unsupervised the task to be learned is some type of association of the input data X with each other (i.e clustering).

Feature engineering

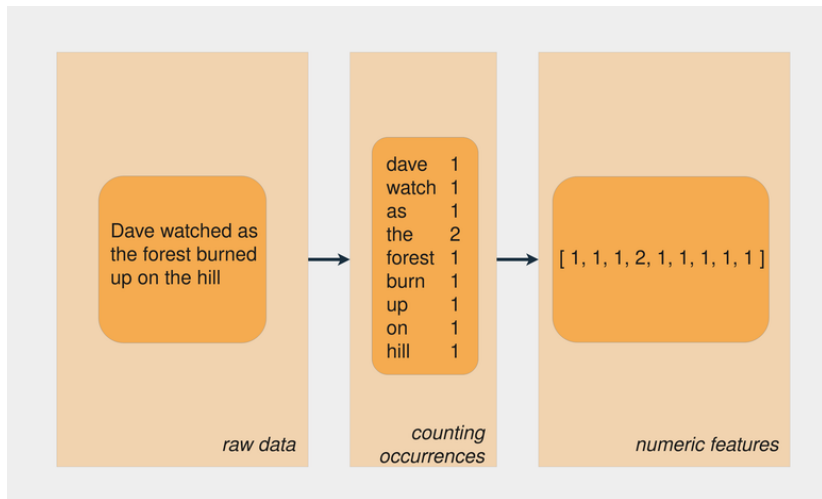


Figure 3: Feature engineering example.


Neural networks

- In the ML example above, we use lemmatization to eliminate inflection from words, generating a simpler model.
- In neural networks, **we leave the network itself** to find out the most important features in the data, without relying on feature engineering techniques.

Neural Networks **are based on** a collection of connected units (neurons), which, just like the synapses in a brain, can transmit a signal to other neurons, so that, acting like interconnected brain cells, they can learn and make decisions in a more human-like manner.

ML vs DL

Machine Learning



Deep Learning




Figure 4: ML vs DL. Source BISMART.

Neural networks




Figure 5: Schema of a 2 layer neural network.

Throwing darts and improving by learning




Figure 6: Steps for trying to hit the center of a dartboard.

Artificial Neural Networks (ANN)

- ANN are formed by layers of neurons.
- Each layer has a certain number of neurons.
- The input data X enters through the first layer and through mathematical operations the input values are transformed into output values y' .
- The goal of the network is to modify the mathematical operations through some parameters (weights and biases, W and b) to minimize the difference between y and y' .
- Deep Learning refers to all those models that use ANN of any type with multiple layers.

Perceptron

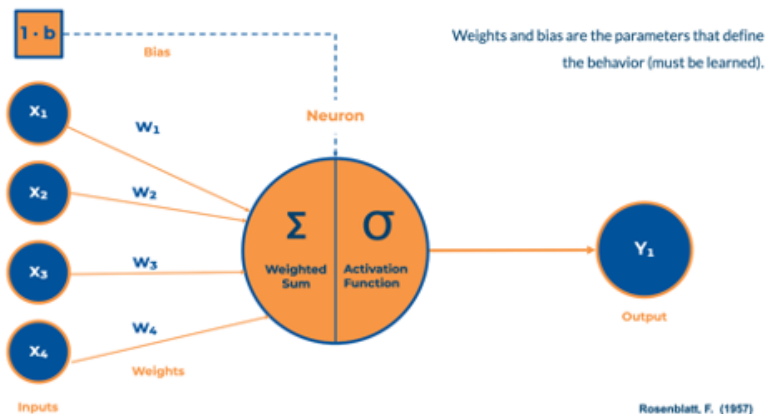


Figure 7: $w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 \leq \text{threshold} \Rightarrow Y = 0$. We can introduce bias and a given function: $\sigma(\mathbf{w} \cdot \mathbf{x} + b) = y'$. Adapted from [1].

Multilayer perceptron




Figure 8: A multilayer perceptron with 5 hidden units.

Multilayer perceptron: some maths

$$w_{ij}^L, \begin{cases} L = \text{layer number} \\ i = \text{neuron from previous layer} \\ j = \text{neuron from following layer} \end{cases}$$

$$\sigma_1 \left(\mathbf{W}^{(1)T} \cdot \mathbf{X} + \mathbf{b}^{(1)} \right) = \sigma_1 \left(\begin{bmatrix} w_{11}^{(1)} & w_{12}^{(1)} & w_{13}^{(1)} & w_{14}^{(1)} \\ w_{21}^{(1)} & w_{22}^{(1)} & w_{23}^{(1)} & w_{24}^{(1)} \\ w_{31}^{(1)} & w_{32}^{(1)} & w_{33}^{(1)} & w_{34}^{(1)} \\ w_{41}^{(1)} & w_{42}^{(1)} & w_{43}^{(1)} & w_{44}^{(1)} \\ w_{51}^{(1)} & w_{52}^{(1)} & w_{53}^{(1)} & w_{54}^{(1)} \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} + \begin{bmatrix} b_1^{(1)} \\ b_2^{(1)} \\ b_3^{(1)} \\ b_4^{(1)} \\ b_5^{(1)} \end{bmatrix} \right) = \mathbf{H}^{(1)}$$

$$\sigma_2 \left(\mathbf{W}^{(2)T} \cdot \mathbf{H}^{(1)} + \mathbf{b}^{(2)} \right) = \sigma_2 \left(\begin{bmatrix} w_{11}^{(2)} & w_{12}^{(2)} & w_{13}^{(2)} & w_{14}^{(2)} & w_{15}^{(2)} \\ w_{21}^{(2)} & w_{22}^{(2)} & w_{23}^{(2)} & w_{24}^{(2)} & w_{25}^{(2)} \\ w_{31}^{(2)} & w_{32}^{(2)} & w_{33}^{(2)} & w_{34}^{(2)} & w_{35}^{(2)} \end{bmatrix} \cdot \begin{bmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \\ h_5 \end{bmatrix} + \begin{bmatrix} b_1^{(2)} \\ b_2^{(2)} \\ b_3^{(2)} \end{bmatrix} \right) = \mathbf{Y}^{\text{out}}$$







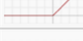


Multilayer perceptron: some details

- The MLP is the simplest neural network
- Weights connect neurons from an inner to an outer layer
- Operations in each layer are implemented using matrices and vectors
- The **forward pass** of the last layer produces the output.
- The MLP can be interpreted as a multivariate function with W and b as the parameters, X as input and Y^{out} as output.
- **Hiperparameters** are additional parameters that control other aspects: number of layers, number of neurons per layer, activation functions, ...

Activation functions

Activation functions:

- The **sigmoid** function is used in the final layer in classification models.
- **ReLU** is used in inner layers, as they produce non-linearity.
- All are differentiable.

Identity		$f(x) = x$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$
Tanh		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$
ArcTan		$f(x) = \tan^{-1}(x)$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU) [2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU) [3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$

An example of a 2 layer Neural Network

Example of classification (for categories "1" and "0") with a 2 layers NN.

- In the first layer we use a linear regression approach.
- The second layer includes the non-linearity through the use of a **sigmoid function** that decides whether the prediction is 1 or 0, following the Bernoulli distribution.




Figure 9: Training a two layer neural network.

The logistic function




Figure 10: A binary classification problem. The class-conditional densities are Gaussians with unit σ^2 . The posterior probability is the logistic function $y = 1/(1 + \exp\{-2x\})$.

Exercise 1

Loss/Cost function in a simple 2 layer NN Check the [provided code](#) for this session and expand it to:

- 1 Generate a collection of 100 synthetic data points from a two Gaussian distribution.
- 2 Assign the data points a category "0" or "1" depending on the Gaussian function they are obtained from.
- 3 Use the provided 2 layers NN to evaluate the expected value of the loss function with the provided weights.
- 4 Can you guess the weights that would improve the prediction (reducing the value of the error)?

Backpropagation: weights

Chain Rule in Backpropagation In each **backward pass**, you compute the partial derivatives of each function, substitute the variables by their values, and finally multiply everything.

Notice that, for the sigmoid function

$$f(x) = \frac{1}{1+e^{-x}}, f'(x) = 1 - f(x).$$

```
error_dweights =
error_dprediction
* dprediction_dlayer1
* dlayer1_dweights
```




Figure 11: Chain rule: derivative of the error with respect to **weights**.

Backpropagation: bias

Chain rule applied to the derivative of the bias error in the 2 layer example.

```

error_dbias =
error_dprediction
* dprediction_dlayer1
* dlayer1_dbias
  
```

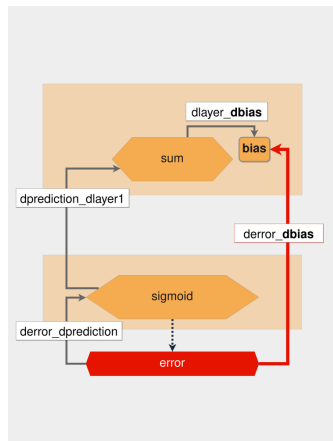


Figure 12: Chain rule: derivative of the error with respect to bias.



F. Rosenblatt.

The perceptron: A probabilistic model for information storage and organization in the brain.

Psychological Review, 65(6):386–408, 1958.

Place: US Publisher: American Psychological Association.



Michael I Jordan.

Why the logistic function? A tutorial discussion on probabilities and neural networks.

Computational Cognitive Science Technical Report 9503, 1995.